



Supervised Classification with daltoolbox

A progressive learning roadmap: from simple baselines to advanced models and hyperparameter selection — all within the unified experimental framework of the daltoolbox package.

R · DALTOOLBOX

DATA MINING

Eduardo Ogasawara

eduardo.ogasawara@cefet-rj.br
<https://eic.cefet-rj.br/~eogasawara>

Supervised vs. Unsupervised Learning

Supervised Learning

Supervised learning is a predictive technique that trains models on labeled data. By learning a mathematical mapping from inputs to outputs, the model can accurately predict values or categories for previously unseen data.

- Relies on labeled data with known input-output pairs
- Learns relationships between features and target outcomes
- Applied to prediction and classification problems

Unsupervised Learning

Unsupervised learning operates on unlabeled data, uncovering hidden structures and patterns without explicit guidance. It can identify natural groupings within the data and serve as a foundation for downstream supervised learning tasks.

- Works with unlabeled, unannotated data
- Discovers inherent structures and patterns
- Can support subsequent prediction tasks through feature learning

Prediction Problems: Classification vs. Regression

Classification

Predicts categorical class labels — whether discrete or nominal. The model is built from a training set and its class attribute values, then applied to classify new, previously unseen data into one of the predefined categories.

Regression

Models continuous-valued functions, enabling the prediction of unknown or missing numerical values. Unlike classification, regression operates on continuous numeric variables rather than discrete categories, making it ideal for forecasting quantities.

Typical Applications



Credit Approval

Determining whether a customer should be approved for a loan or credit card based on their financial history and credit profile.



Medical Diagnosis

Identifying whether a tumor is malignant or benign by analyzing clinical features and patient data through a trained classification model.



Fraud Detection

Classifying financial transactions as fraudulent or legitimate to protect banking and payment systems from malicious activity.



Web Categorization

Assigning appropriate categories to web pages to facilitate content organization, search indexing, and information retrieval.

Classification – A Two-Step Process

1

Step 1: Model Construction

The first phase involves describing a set of predetermined classes. Each tuple or sample belongs to a predefined class, determined by its class label attribute. The collection of tuples used to build the model is known as the **training set**.

- The model can be represented as classification rules, decision trees, or mathematical formulas
- Training data is used to learn patterns that generalize to unseen examples

2


Step 2: Model Application

Once constructed, the model is applied to classify future or previously unseen objects. Its accuracy is estimated by comparing the known class labels of test samples against the model's predicted results.

- **Accuracy rate:** the percentage of test set samples correctly classified by the model
- The test set must be independent of the training set to avoid overfitting
- If the accuracy is deemed acceptable, the model can be deployed to classify new, real-world data

❏ **Important Note:** During training, if it is necessary to select among different candidate models, a portion of the training set should be set aside for evaluation — this subset is referred to as the **validation set**.

Training the Model



- **Classification Algorithms**
Process the training data to identify underlying patterns
- **Classifier (Model)**
The resulting structure capable of making predictions

Example of a Learned Rule:


```
IF rank = 'professor'  
  OR years > 6  
THEN tenured = 'yes'
```

The training set contains labeled examples that the algorithm uses to learn patterns and relationships between features and classes.

This example illustrates how a model can learn logical rules from the training set, constructing conditions that determine the classification of new, unseen cases.

Reference: J. Han, J. Pei, and H. Tong, Data Mining: Concepts and Techniques, 4th edition. Cambridge, MA: Morgan Kaufmann, 2022.

Using the Model for Prediction



Input Data

Name: Merlisa

Role: Associate Prof

Years of Experience: 7

The model receives a new case with known features, but without an assigned class label.

Classification

The model applies the rules learned during training to determine the appropriate class.

Prediction: Tenured?

Based on the rule *IF years > 6*, the model classifies Merlisa as a candidate for tenure.

This process illustrates how the model generalizes the knowledge acquired during training to make predictions on new, unseen cases — applying the learned rules and patterns to produce accurate classifications.

Learning Roadmap

The examples are organized into three progressive groups. Study them in this order to build a solid understanding before advancing to more complex models.

01

Foundations

Majority baseline and Decision Tree — learn the complete cycle: data splitting, training, prediction, and evaluation.

02

Model Families

KNN, Naive Bayes, Random Forest, SVM, and MLP — explore different supervised learning paradigms and their classification strategies.

03

Model Selection

Hyperparameter search with `cla_tune` — fine-tune your models without breaking the experimental structure you have already learned.

Summary: Supervised Classification with daltoolbox

Classification as a Two-Step Process

Supervised classification follows a two-phase cycle: **Model Construction** — learning rules from labeled training data — followed by **Model Application** — predicting classes for unseen cases. The daltoolbox enforces this discipline consistently through `fit()`, `predict()`, and `evaluate()`.

Progressive Model Complexity

The learning roadmap builds from simple to sophisticated: start with the **Majority Baseline** and **Decision Tree** (Foundations), then explore **KNN** (distance-based), **Naive Bayes** (probabilistic), **Random Forest** (ensemble), **SVM** (kernel-based), and **MLP** (neural networks) in the Model Families group.

Hyperparameter Tuning with `cla_tune`

The Model Selection stage fine-tunes learned models using `cla_tune` within the same unified daltoolbox experimental framework. No structural changes are needed — the same split → fit → predict → evaluate workflow applies, ensuring reproducible and comparable results across all experiments.

References & Additional Sources

Academic references and online resources to explore the DAL Toolbox and Experiment Lines further.

Online Resources

CRAN Package – daltoolbox

cran.r-project.org/web/packages/daltoolbox

DAL Toolbox – GitHub

github.com/cefet-rj-dal/daltoolbox

DAL Toolbox – Documentation

cefet-rj-dal.github.io/daltoolbox

Official documentation, tutorials, and vignettes for the daltoolbox package.

Key Academic References

- Ogasawara, E., Sa, A., Castro, A., Santos, C., Salles, D., Lima, J., Tavares, L., Salles, R. & Saidy, V. (2026). daltoolbox: Leveraging Experiment Lines to Data Analytics. R package version 1.3.767. *CRAN*.
- Ogasawara, E., Paulino, C., Murta, L., Werner, C. & Mattoso, M. (2009). Experiment Line: Software Reuse in Scientific Workflows. *SSDBM 2009, LNCS 5566*, pp. 264–272. *Springer*.
- Han, J., Pei, J. & Tong, H. (2022). Data Mining: Concepts and Techniques, 4th ed. *Morgan Kaufmann*. (decision trees, Naive Bayes, SVM)
- Breiman, L. (2001). Random Forests. *Machine Learning*, 45(1), 5–32. *Springer*.
- Cover, T. & Hart, P. (1967). Nearest neighbor pattern classification. *IEEE Transactions on Information Theory*, 13(1), 21–27.
- Cortes, C. & Vapnik, V. (1995). Support-vector networks. *Machine Learning*, 20(3), 273–297. *Springer*.
- Rumelhart, D. E., Hinton, G. E. & Williams, R. J. (1986). Learning representations by back-propagating errors. *Nature*, 323, 533–536.