



Extending a Machine Learning Framework in R: Custom Learners Explained

Custom Examples in daltoolbox

The **daltoolbox** keeps the workflow stable while the user connects a new backend through a small integration contract. This series of examples organizes extensions by role: preprocessing, supervised learning, and unsupervised components.

01

Data Preparation

Custom transformations before modeling

02

Supervised Learning

Custom classifiers and regressors

03

Unsupervised Learning

Clustering with external backends

Custom Transformation

SMOTE (Synthetic Minority Over-sampling Technique) is a balancing algorithm that generates synthetic samples for minority classes through **interpolation** between existing minority instances and their k -nearest neighbors in the feature space — rather than simply duplicating samples. This avoids overfitting and produces more realistic decision boundaries.

① Select Sample

A minority class instance is randomly chosen as the starting point.

② Find k -Neighbors

The k nearest neighbors of that instance are identified in the feature space.

③ Interpolate Synthetic Point

A new point is generated along the line segment between the sample and a randomly selected neighbor: $x_{\text{new}} = x + \lambda \cdot (x_{\text{nn}} - x)$, with $\lambda \in [0, 1]$.

Integration Contract

- Constructor: `smote_custom(attribute)`
- Method: `transform.smote_custom(obj, data, ...)`
- External backend: `smotefamily::SMOTE`

Result on Imbalanced Iris

Before: setosa=50, versicolor=21, virginica=11

After: setosa=50, versicolor=42, virginica=44

- ① Chawla, N. V., Bowyer, K. W., Hall, L. O., & Kegelmeyer, W. P. (2002). SMOTE: Synthetic Minority Over-sampling Technique. *Journal of Artificial Intelligence Research*, 16, 321–357.
- He, H., & Garcia, E. A. (2009). Learning from Imbalanced Data. *IEEE Transactions on Knowledge and Data Engineering*, 21(9), 1263–1284.

Custom Data Preparation

```
library(daltoolbox)

smote_custom <- function(attribute) {
  obj <- daltoolbox::dal_transform()
  obj$attribute <- attribute
  class(obj) <- append("smote_custom", class(obj))
  obj
}

transform.smote_custom <- function(
  obj, data, ...) {
  data <- data.frame(data, check.names = FALSE)
  j <- match(obj$attribute, colnames(data))
  x <- sort(table(data[, obj$attribute]))
  # oversampling loop per class
  # using smotefamily::SMOTE
  ...
}
```

```
iris_imb <- datasets::iris[
  c(1:50, 51:71, 101:111), ]
bal <- smote_custom("Species")
iris_bal <- transform(bal, iris_imb)
table(iris_bal$Species)
# setosa:50 versicolor:42 virginica:44
```

Step by Step

- **Constructor:** calls `dal_transform()`, stores the target attribute and adds the custom class.
- **S3 Method:** implements `transform.smote_custom()` — the only method required by the contract.
- **External backend:** delegates oversampling to `smotefamily::SMOTE`, outside the `daltoolbox` namespace.
- **Result:** balanced classes (50 / 42 / 44) without altering the rest of the pipeline.

📌 **Key concept:** The integration contract is minimal — just a constructor + S3 method. The analytical workflow remains unchanged.

Custom Classification

A **Multilayer Perceptron (MLP)** is a feedforward artificial neural network composed of an input layer, one or more hidden layers, and an output layer. Each neuron applies a **non-linear activation function** (e.g., sigmoid or ReLU) to a weighted sum of its inputs. Training uses **backpropagation** with gradient descent to minimize a loss function. MLPs are universal function approximators (Cybenko, 1989) and are widely used in classification tasks.

Input Layer

Receives the feature vector (e.g., 4 attributes from Iris) and propagates them to the next layer.

Output Layer

Produces per-class probabilities via **softmax**. The class with the highest probability is the final prediction.

Hidden Layer(s)

Applies synaptic weights + sigmoid activation function: $\sigma(z) = \frac{1}{1+e^{-z}}$. The `size` hyperparameter defines the number of neurons.

Backpropagation

Adjusts weights by minimizing the **cross-entropy loss** via gradient descent. Controlled by `learn_rate` and `maxit`.

Hyperparameters

- `size` — neurons in the hidden layer (default: 5)
- `learn_rate` — learning rate (default: 0.2)
- `maxit` — maximum iterations (default: 100)

Training (Iris)

Accuracy: **0.9583**
Precision, Recall, F1: **1.0**

Test (Iris)

Accuracy: **0.9333**
Precision, Recall, F1: **1.0**

📖 **References:** Rumelhart, D. E., Hinton, G. E., & Williams, R. J. (1986). Learning representations by back-propagating errors. *Nature*, 323, 533–536. — Cybenko, G. (1989). Approximation by superpositions of a sigmoidal function. *Mathematics of Control, Signals and Systems*, 2(4), 303–314.

Custom Classifier

```
library(daltoolbox)

cla_rsnnns_custom <- function(
  attribute, slevels, size = 5, learn_rate = 0.2, maxit = 100) {
  obj <- daltoolbox::classification(attribute, slevels)
  obj$size <- size
  obj$learn_rate <- learn_rate
  obj$maxit <- maxit
  class(obj) <- append("cla_rsnnns_custom", class(obj))
  obj
}

fit.cla_rsnnns_custom <- function(
  obj, data, ...) {
  x <- as.matrix(data[, obj$x, drop = FALSE])
  y <- daltoolbox::adjust_class_label(data[, obj$attribute])
  obj$model <- RSNNs::mlp(x = x, y = y, size = obj$size,
    learnFuncParams = c(obj$learn_rate), maxit = obj$maxit)
  obj
}
```

```
predict.cla_rsnnns_custom <- function(
  object, x, ...) {
  x <- as.matrix(x[, object$x, drop = FALSE])
  prediction <- predict(object$model, x)
  colnames(prediction) <- object$slelevels
  as.data.frame(prediction)
}

# Train / Test
model <- cla_rsnnns_custom("Species", slelevels,
  size = 5, learn_rate = 0.1, maxit = 150)
model <- fit(model, iris_train)
# accuracy train: 0.9583
# accuracy test: 0.9333
```

❏ **Key concept:** Constructor + fit() + predict() is the complete contract. The algorithm complexity stays in the backend; the integration remains small and predictable.

Step by Step

- **Constructor:** inherits from `classification()` and stores hyperparameters (`size`, `learn_rate`, `maxit`).
- **fit():** prepares X and Y matrices, trains the MLP via `RSNNs::mlp` and returns the updated object.
- **predict():** returns a `data.frame` with columns named by levels — the format expected by `evaluate()`.
- **Downstream:** sampling and evaluation are identical to native `daltoolbox` examples.

Custom Regression

To adapt an MLP for regression, the output layer uses **linear activation** (`linOut = TRUE` in RSNNS), allowing the network to predict continuous values without restricting the output to the `[0,1]` interval. The loss function changes from **cross-entropy** (classification) to **Mean Squared Error (MSE)**. The hidden layer architecture remains identical to the classification MLP — only the output contract and loss function differ.

Output Activation

Classification: softmax

Regression: linear (identity)

Loss Function

Classification: cross-entropy

Regression: MSE (mean squared error)

Integration Contract

- **Constructor:** `reg_rsnns_custom(attribute, size, learn_rate, maxit)`
- **fit():** trains MLP with `linOut = TRUE` via `RSNNS::mlp`
- **predict():** returns a numeric vector (no class columns)

Return Type

Classification: probabilities per class

Regression: continuous numeric vector

RSNNS Parameter

Classification: `linOut = FALSE`

Regression: `linOut = TRUE`

Results – Boston Dataset

Train: MSE = 108.64 · SMAPE = 0.369 · $R^2 = -0.207$

Test: MSE = 95.48 · SMAPE = 0.378 · $R^2 = -0.587$

The negative R^2 values indicate that the model, without hyperparameter tuning, performs below a simple baseline — expected for a first standard run.

- ❏ **References:** Hornik, K., Stinchcombe, M., and White, H. (1989). Multilayer feedforward networks are universal approximators. *Neural Networks*, 2(5), 359–366. — Bergmeir, C., and Benitez, J. M. (2012). Neural Networks in R Using the Stuttgart Neural Network Simulator: RSNNS. *Journal of Statistical Software*, 46(7), 1–26.

Custom Regressor

```
library(daltoolbox)
reg_rsnns_custom <- function(attribute, size = 5, learn_rate = 0.1, maxit = 200) {
  obj <- daltoolbox::regression(attribute)
  obj$size <- size
  obj$learn_rate <- learn_rate
  obj$maxit <- maxit
  class(obj) <- append("reg_rsnns_custom", class(obj))
  obj
}

fit.reg_rsnns_custom <- function(
  obj, data, ...) {
  x <- as.matrix(data[, obj$x, drop = FALSE])
  y <- as.matrix(data[, obj$attribute, drop = FALSE])
  obj$model <- RSNNS::mlp(
    x = x, y = y, size = obj$size,
    learnFuncParams = c(obj$learn_rate),
    maxit = obj$maxit, linOut = TRUE) # linear output!
  obj
}
```

```
predict.reg_rsnns_custom <- function(
  object, x, ...) {
  x <- as.matrix(x[, object$x, drop = FALSE])
  as.numeric(predict(object$model, x))
}

# Boston Dataset (MASS)
model <- reg_rsnns_custom("medv", size = 6, learn_rate = 0.05, maxit = 200)
model <- fit(model, boston_train)
# MSE train: 108.64
# MSE test: 95.48
```

Step by Step

- **Constructor:** inherits from `regression()`; hyperparameters are stored in the object.
- **fit():** uses `linOut = TRUE` to enable continuous linear output — the essential difference from the classifier.
- **predict():** returns a simple numeric vector, compatible with `evaluate()`.
- **Separation of concerns:** the algorithm can be complex; the integration contract remains identical to that of the classifier.

📌 **Key concept:** The only structural difference between a custom classifier and a custom regressor is `linOut = TRUE` and the return type of `predict()`.

Custom Clustering

AGNES (Agglomerative Nesting) is a *bottom-up* hierarchical clustering algorithm. It starts with each observation as its own cluster and iteratively merges the two most similar clusters until only a single cluster remains. The result is a **dendrogram** — a tree structure that can be cut at any level to produce k clusters. The linkage method (e.g., Ward) determines how the distance between clusters is computed.

1 • Initialize

Each observation forms its own cluster — n clusters for n points.

2 • Compute Distances

Builds the distance matrix using the chosen metric (e.g., Euclidean).

3 • Merge

Joins the two closest clusters according to the linkage criterion — Ward minimizes intra-cluster variance.

4 • Repeat

Recomputes distances and repeats the merge until all observations belong to a single cluster.

5 • Cut

Applies `cutree(k)` to extract k flat clusters from the dendrogram.

Integration Contract

- **Constructor:** inherits from `clustering()`; stores `k` and the linkage method.
- **fit():** calls `cluster::agnes()` and stores the learned hierarchical structure.
- **cluster():** applies `cutree()` and returns a vector of integer labels, compatible with `evaluate()`.

Results — Iris ($k=3$, Ward)

- Cluster 1: **49** observations
- Cluster 2: **30** observations
- Cluster 3: **71** observations

Entropy per cluster: $CE_1=0.000 \cdot CE_2=0.561 \cdot CE_3=0.909$
Total entropy: **0.542**

📖 **References:** Kaufman, L., and Rousseeuw, P. J. (1990). *Finding Groups in Data: An Introduction to Cluster Analysis*. Wiley. — Ward, J. H. (1963). Hierarchical Grouping to Optimize an Objective Function. *Journal of the American Statistical Association*, 58(301), 236–244.

Custom Clusterer

```
library(daltoolbox)

cluster_agnes_custom <- function(
  k = 3, method = "ward", metric = "euclidean", scale = TRUE) {
  obj <- daltoolbox::clusterer()
  obj$k <- k
  obj$method <- method
  obj$metric <- metric
  obj$scale <- scale
  class(obj) <- append("cluster_agnes_custom", class(obj))
  obj
}

fit.cluster_agnes_custom <- function(
  obj, data, ...) {
  x <- as.matrix(data)
  storage.mode(x) <- "double"
  if (isTRUE(obj$scale)) x <- scale(x)
  obj$model <- cluster::agnes(
    x, diss = FALSE,
    metric = obj$metric,
    method = obj$method)
  obj
}
```

```
cluster.cluster_agnes_custom <- function(obj, data, ...) {
  x <- as.matrix(data)
  storage.mode(x) <- "double"
  clu <- stats::cutree(as.hclust(obj$model), k = obj$k)
  # calcula métrica intra-cluster
  ...
  attr(clu, "metric") <- dist
  clu
}

# Iris — 4 features, k=3
model <- cluster_agnes_custom(k = 3, method = "ward")
model <- fit(model, iris[, 1:4])
clu <- cluster(model, iris[, 1:4])
# clu: 1→49 2→30 3→71
eval <- evaluate(model, clu, iris$Species)
# clustering_entropy: 0.542
```


Step by Step

- **Constructor:** inherits from `clusterer()`; stores `k`, linkage method, and distance metric.
- **fit():** scales the data (optional) and trains the agglomerative hierarchical model via `cluster::agnes`.
- **cluster():** cuts the dendrogram with `cutree()` and computes the intra-cluster distance as a quality metric.
- **evaluate():** returns per-cluster entropy and global entropy — same standard as native clusterers.

❏ **Key concept:** The contract for clustering is constructor + `fit()` + `cluster()`. The workflow remains stable even with an external algorithm.

The Framework's Unifying Principle

Regardless of the extension type — transformation, classification, regression, or clustering — daltoolbox maintains the same small and uniform integration contract.



❏ **Conclusion:** The mining method can be sophisticated, but the integration contract remains small, stable, and reusable at any point in the Experiment Line.

Acknowledgements



Data Analytics Lab

dal.global@grupo.cefet-rj.br | <https://eic.cefet-rj.br/dal>

